

DISHA LAMBA

jb.dishalamba@gmail.com | dldisha.github.io | linkedin.com/in/dldisha/

PROFESSIONAL SUMMARY

Data Scientist with 4+ years of experience building production-grade machine learning systems, data pipelines, and analytics solutions on large-scale datasets. Experienced in developing and deploying LLM applications, including fine-tuning foundation ML models, RAG pipelines, prompt engineering, and evaluation frameworks. Skilled in Python, SQL, GCP, with a proven ability to take AI products from experimentation to production while collaborating cross-functionally to deliver measurable business outcomes.

WORK EXPERIENCE

CVS Health

New York, NY

Data Scientist

Aug 2024 – May 2026

- Architected and shipped a **production LLM pipeline** that extracts structured clinical insights from unstructured discharge notes for a **1.2M**-member population, reducing clinician **review time by 80%** through schema-guided prompting (**Pydantic**) and hallucination detection, and enabling downstream reporting, care management, and operational decision-making.
- Built a **chronic condition risk-prediction model** combining **RAG-based fact extraction** with **structured LLM prompting** over claims and clinical data to identify high-risk patient cohorts projected to drive **~\$100M in cost avoidance**.
- Led **controlled experimentation** comparing model outputs against medical-director review across patient subgroups, using **significance testing (precision, recall)** and **Bayesian estimation** to quantify reliability and uncertainty before rollout.
- Built automated data quality validation and production monitoring pipelines to ensure schema integrity, data completeness, and trusted analytical datasets by proactively detecting anomalies before impacting downstream reporting.
- Consolidated **20M+ records** across disparate clinical sources into a unified, well-modeled BigQuery dataset using **complex SQL queries (joins, window functions, CTEs)**, eliminating **70% of ad hoc data requests** from product and engineering teams.
- Partnered with Product stakeholders to define and standardize **20+ product, engagement, and operational KPIs**, building **Tableau dashboards** to track feature adoption, user engagement, retention, and support executive decision-making.
- Collaborated closely with Product, Engineering, and clinical leadership to translate ambiguous business questions into analytical solutions, defining success metrics, evaluating outcomes, and prioritizing roadmap initiatives using data.
- Mentored **3 junior associates** day-to-day across task planning, code reviews, and feedback to keep delivery on track.

Mindgrasp

Columbia, MD

Product Data Scientist

Jan 2024 – April 2024

- Designed and shipped an AI-powered Math Q&A feature end-to-end by integrating **Wolfram Alpha with an AWS Lambda** backend and **React front end**, achieving **38% adoption** among existing users within 2 months of launch.
- Ran the trial-to-paid conversion workstream, analyzing the signup-to-paid funnel in **Mixpanel** (DAU/MAU, retention, cohort conversion) and A/B testing paywall placement to lift **trial-to-paid conversion 12%**.

Subconscious AI

New York, NY

Machine Learning Engineer

May 2023 – Jan 2024

- Built, trained, and integrated discrete choice causal models to **simulate customer populations as digital twins**, achieving **93% predictive accuracy** against real human behavioral benchmarks validated **across 200+ published studies**, enabling clients across pharma, CPG, and consulting to pressure-test pricing, messaging, and product decisions before market launch.

SKILLS

Languages/Frameworks: Python, SQL, PostgreSQL, PySpark, Pandas, NumPy, XGBoost, Matplotlib, Scikit-learn, Keras, HTML/CSS

Machine Learning & AI: LLMs, RAG, Prompt Engineering, NLP, Classification, Regression, Clustering, Model Evaluation

Statistics & Experimentation: A/B Testing, Hypothesis Testing, Time-Series Analysis, Statistical Significance, Cohort Analysis

Cloud & Other Tools: AWS, GCP, Snowflake, dbt, Mixpanel, Segment, Tableau, Power BI, Git, Jira, Figma, Confluence

PROJECTS

Recommendation System using MovieLens Dataset (SQL, Spark)

- Built a scalable movie recommendation system on MovieLens 32M dataset using PySpark & SparkSQL, applying ALS algorithm for latent factor modelling, benchmarking against user-based and popularity-based approaches to maximise accuracy.

EDUCATION

New York University, Master of Science in Computer Engineering, May 2023

Guru Gobind Singh Indraprastha University, Bachelor of Technology in Information Technology, Sept 2020

LEADERSHIP & RECOGNITION

- Speaker**, AICamp'24 NYC, Data Summit'25 Boston, Selected Speaker at Women Who Code'26
- Best Medical Hack** at [MHacks 13 hackathon](#); developed ReactJS web app to match COVID-19 patients with plasma donors.